

Predicting Patient Survival: A Classification Approach

Rahul Kumar Sahoo
Politecnico di Torino
Student id: s316411
s316411@studenti.polito.it

Nithyakalyani Karthikeyan
Politecnico di Torino
Student id: s310149
s310149@studenti.polito.it

Abstract—The project aims to develop a classification model that can correctly forecast a patient’s binary outcome of “death” in relation to their likelihood of surviving within a given amount of time. After validating a number of models, we ultimately decided to use the random forest classifier to forecast our results. The model obtained an F1score of 0.748.

I. PROBLEM OVERVIEW

A. Aim

The primary goal of this project is to develop a machine learning model that accurately predicts the binary outcome “death” for these patients, which signifies whether a patient survives or dies within a specified time frame.

B. Challenges

The primary challenge in this project was handling the significant amount of missing data within the dataset. Specifically, 4 to 5 columns had more than 40% null values, and several other columns had between 20% to 30% null values. Given the critical nature of this medical dataset, every column potentially contributes valuable information towards predicting patient survival. For instance, abnormal results in columns like blood pH or glucose levels can be indicative of severe medical conditions and can influence the survival prediction. Therefore, it was crucial to decide on an appropriate method to handle these missing values without compromising the integrity of the data.

Initially, we considered several imputation techniques to address the missing values. One of the methods attempted was MICE (Multiple Imputation by Chained Equations) [7], which iteratively fills in each column’s missing values based on the other columns. Although MICE is a sophisticated technique, it did not produce satisfactory results in our case. The complexity and iterative nature of MICE imputation posed challenges in terms of computational resources and convergence stability.

Given the unsatisfactory performance of MICE, we explored simpler imputation methods. Ultimately, we opted for mean imputation, which involves replacing missing values with the mean of the respective column. This method, while straightforward, provided a decent solution and helped maintain the overall structure and distribution of the data.

Another challenge was ensuring that the chosen imputation method did not introduce bias or distort the relationships between variables. The integrity of the data is paramount in medical datasets, as incorrect imputation could lead to

inaccurate predictions and potentially harmful decisions in a clinical setting. Therefore, we meticulously evaluated the impact of mean imputation [8] on the dataset to ensure it was a viable solution.

In addition to handling missing values, we faced the challenge of selecting the right features and encoding categorical variables. We performed one-hot encoding for non-hierarchical categorical variables such as ‘dzclass,’ ‘ca,’ and ‘dnr,’ transforming them into a suitable format for machine learning algorithms. This preprocessing step was essential to preserve the categorical information without introducing ordinal relationships that did not exist.

Distribution analysis revealed uneven distributions in several columns, including ‘prg6m,’ ‘prg2m,’ ‘dzgroup,’ and ‘income.’ Handling these uneven distributions was critical to ensure that the model did not become biased towards certain values. We used various statistical techniques to understand and address these distribution issues, ensuring a more balanced and representative dataset for training the model.

Overall, the main challenges revolved around handling missing data, selecting appropriate imputation methods, encoding categorical variables, and addressing distribution imbalances. These preprocessing steps were crucial in developing a reliable and accurate machine learning model for predicting patient survival. The careful consideration and resolution of these challenges underscore the importance of data integrity and preprocessing in the development of robust predictive models in the medical domain.

C. Development dataset

The development data consists of 7284 rows and 44 columns. Each row concerns hospitalized patient records who met the inclusion and exclusion criteria for nine disease categories: acute respiratory failure, chronic obstructive pulmonary disease, congestive heart failure, liver disease, coma, colon cancer, lung cancer, multiple organ system failure with malignancy, and multiple organ system failure with sepsis which were recorded throughout 1989-1991 and 1992-1994 in America.

D. Evaluation dataset

The evaluation data consists of 1821 rows and 43 columns.

We will need to use the development set to build a classification model to correctly label the points in the

- kNN [6]: KNearest Neighbors *KNN* is a simple, yet effective supervised learning algorithm used for classification and regression tasks. It operates by identifying the k -nearest data points to a given input and making predictions based on the majority class (for classification) or the average value (for regression) of these neighbors. KNN is non-parametric and instance-based, meaning it makes decisions based on the entire training dataset without assuming an underlying distribution. The algorithm is particularly useful for datasets with a clear cluster structure. However, KNN can be computationally intensive for large datasets and sensitive to the choice of k and the distance metric used. It performs well in low-dimensional spaces but may struggle with high-dimensional data due to the curse of dimensionality.
- Naive Bayes [5] :Naive Bayes is a collection of supervised learning algorithms based on Bayes' theorem, with the simplifying assumption that features are conditionally independent given the class variable. Despite this "naive" assumption, Naive Bayes classifiers have demonstrated strong performance in real-world applications such as document classification and spam filtering. These classifiers require minimal training data to estimate parameters and are extremely fast compared to more complex methods. Each feature's distribution is estimated independently, which helps mitigate issues related to high-dimensional data. However, while Naive Bayes is effective as a classifier, its probability estimates are often unreliable and should not be heavily relied upon.




Fig. 2. Model comparison

Among the algorithms, random forest demonstrated superior performance compared to others. In medical datasets, the random forest classifier often performs exceptionally well because each column typically corresponds to a specific organ or physiological parameter. If an organ is failing, it can be a significant predictor of patient mortality. This scenario lends itself well to a tree structure approach, where a decision tree can evaluate conditions sequentially. For example, if glucose

levels are high, the model can then assess the next relevant parameter, continuing this process until it reaches a final prediction of either survival or death. This hierarchical, tree-based learning method enables the model to effectively and efficiently learn from the medical dataset. While decision trees alone perform well in this context, random forests, which are an ensemble of multiple decision trees, tend to outperform single decision trees due to their ability to reduce overfitting and improve predictive accuracy. Thus, the random forest classifier is particularly well-suited for medical datasets, combining the strengths of individual decision trees into a more robust predictive model. Given its performance, we selected the random forest algorithm for our final model.

C. Hyperparameter Tuning

We proceeded to fine-tune the hyperparameters of the random forest model. Despite trying automated methods such as grid search and random search CV, we did not achieve optimal results. Therefore, we manually

tuned the hyperparameters, leading to the following configuration:

- `n_estimators = 30`
- `max_depth = 30`
- `random_state = 42`
- `class_weight = 'balanced'`
- `criterion = 'log_loss'`

This set of hyperparameters provided the best performance, yielding an F1 score of 0.748 on the evaluation dataset.

III. RESULTS

The Random Forest model outperformed other models during comparison, yielding an F1 macro score of 0.748 on the evaluation dataset and 0.749 on the development dataset. The model achieved an accuracy of 78% on the development dataset, which is quite satisfactory. However, as observed in the table below, the model performs better at predicting class 1.0 than class 0.0, as indicated by higher precision, recall, and F1 scores for class 1.0.

Given the imbalance between class 0 and class 1, we attempted to address this by resampling the data. After resampling, the model's F1 macro score on the development dataset increased to 0.86. Unfortunately, this improvement did not translate to the evaluation dataset, where the F1 score decreased to 0.743. Consequently, we decided to keep the original data unchanged and revisit the resampling approach for further refinement.

IV. CLASSIFICATION REPORT FOR DEVELOPMENT DATASET

	precision	report	f1 score	support
0.0	0.67	0.61	0.64	462
1.0	0.83	0.86	0.84	995
accuracy			0.78	1457
macro avg	0.75	0.74	0.74	1457
weighted avg	0.78	0.78	0.78	1457

TABLE I
CLASSIFICATION REPORT




Fig. 3. Confusion matrix for development dataset

V. DISCUSSION

In summary, our preprocessing involved careful handling of missing values, distribution analysis, and one-hot encoding of categorical variables. We selected the random forest classifier based on its initial performance and further improved it through manual hyperparameter tuning. The final model achieved an F1 score of 0.748, demonstrating its efficacy in predicting patient survival.

This comprehensive approach underscores the importance of thoughtful data preprocessing and targeted model tuning, especially in critical applications such as medical predictions.

REFERENCES

- [1] <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>
- [2] https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html
- [3] <https://scikit-learn.org/stable/modules/svm.html>
- [4] <https://scikit-learn.org/stable/modules/tree.html>
- [5] https://scikit-learn.org/stable/modules/naive_bayes.html
- [6] <https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>
- [7] <https://www.machinelearningplus.com/machine-learning/mice-imputation>
- [8] <https://scikit-learn.org/stable/modules/generated/sklearn.impute.SimpleImputer.html>